



SQL Server Architecture

Module 2

Learning Units covered in this Module

- Lesson 1: Introduction to SQL Operating System
- Lesson 2: SQL Server Waits and Queues

Lesson 1: Introduction to SQL Operating System


Objectives

After completing this learning, you will be able to:

- Understand Microsoft SQL Server Architecture Overview.
- Learn SQL Server Operating System (SQLOS) Components.



Inside the Database Engine



Two Main Functions of SQLOS

Management

- Memory Manager
- Process Scheduler
- Synchronization
- I/O
- Support for Non-Uniform Memory Access (NUMA) and Resource Governor

Monitoring

- Resource Monitor
- Deadlock Monitor
- Scheduler Monitor
- Lazy Writer (Buffer Pool management)
- Dynamic Management Views (DMVs)
- Extended Events
- Dedicated Administrator Connection (DAC)

Dynamic Management Views and Functions

Category	Description
sys.dm_exec_%	Execution and connection information
sys.dm_os_%	Operating system related information
sys.dm_tran_%	Transaction management information
sys.dm_io_%	I/O related information
sys.dm_db_%	Database information

Using Dynamic Management Objects (DMOs)

- Must reference using the sys schema
- Two basic types:
 - Real-time state information
 - Historical information

```
SELECT cpu_count, hyperthread_ratio,  
       scheduler_count, scheduler_total_count,  
       affinity_type, affinity_type_desc,  
       softnuma_configuration, softnuma_configuration_desc,  
       socket_count, cores_per_socket, numa_node_count,  
       sql_memory_model, sql_memory_model_desc  
FROM sys.dm_os_sys_info
```


SQL Server Configuration

Processor Configuration Settings And Best Practices

Affinity Mask

- Assigns CPUs for SQL Server use
- Set via `sp_configure` or `Alter Server Configuration`
- Only required in specific scenarios

Max Degree of Parallelism (MAXDOP)

- Maximum number of processors that are used for the execution of a query in a parallel plan. This option determines the number of threads that are used for the query plan operators that perform the work in parallel.

Cost Threshold for Parallelism

- Only queries with a cost that is higher than this value will be considered for parallelism
- Only required when dealing with excessive parallelism

Max Worker Threads

- Number of threads SQL Server can allocate
- Recommended value is 0. SQL Server will dynamically set the Max based on CPUs and CPU architecture

Demonstration

Dynamic Management Views



Questions?



Lesson 2: Waits and Queues

Objectives

After completing this learning, you will be able to:

- Understand SQL Server task scheduling.
- Understand Waits and Queues



Microsoft SQL Server Scheduling Terminology

Batch

- A statement or set of statements submitted to SQL Server by the user (a query), also referred to as a request
- Monitor with `sys.dm_exec_requests`

Task

- A batch will have one or more tasks (aligns with statements)
- Monitor with `sys.dm_os_tasks`

Worker Thread

- Each task will be assigned to a single worker thread for the life of the task
- Monitor with `sys.dm_os_workers`

Hierarchy of Common Terms

```
SELECT *  
FROM sys.dm_exec_sessions;  
-- relevant data:  
-- session_id --> spid, memory address of the worker  
-- task_id --> batch_id, request_id, time slice, etc.  
-- status --> running, suspended, etc.  
-- wait_type --> type of wait, e.g. IO, LCK, etc.  
-- wait_time --> time spent waiting  
-- open_transaction_count --> blocking identification  
-- cpu_time, total_elapsed_time, reads, writes --> telemetry
```

Connection

Session

Request
(Batch)

Task

Worker

Scheduling Types

Non-Preemptive (Cooperative)

- SQL Server manages CPU scheduling for most activity (instead of the operating system).
- SQL Server decides when a thread should wait or get switched out (known as yielding).
- SQL Server developers also programmed some predetermined voluntary yields to avoid starvation of other threads

Preemptive

- Preemption is the act of an operating system temporarily interrupting an executing task.
- Higher priority tasks can preempt lower priority tasks.
- Preemptive mode used in SQL Server for external code calls, CLR with an UNSAFE assemblies, extended stored procedures

SQL Server Task Scheduling

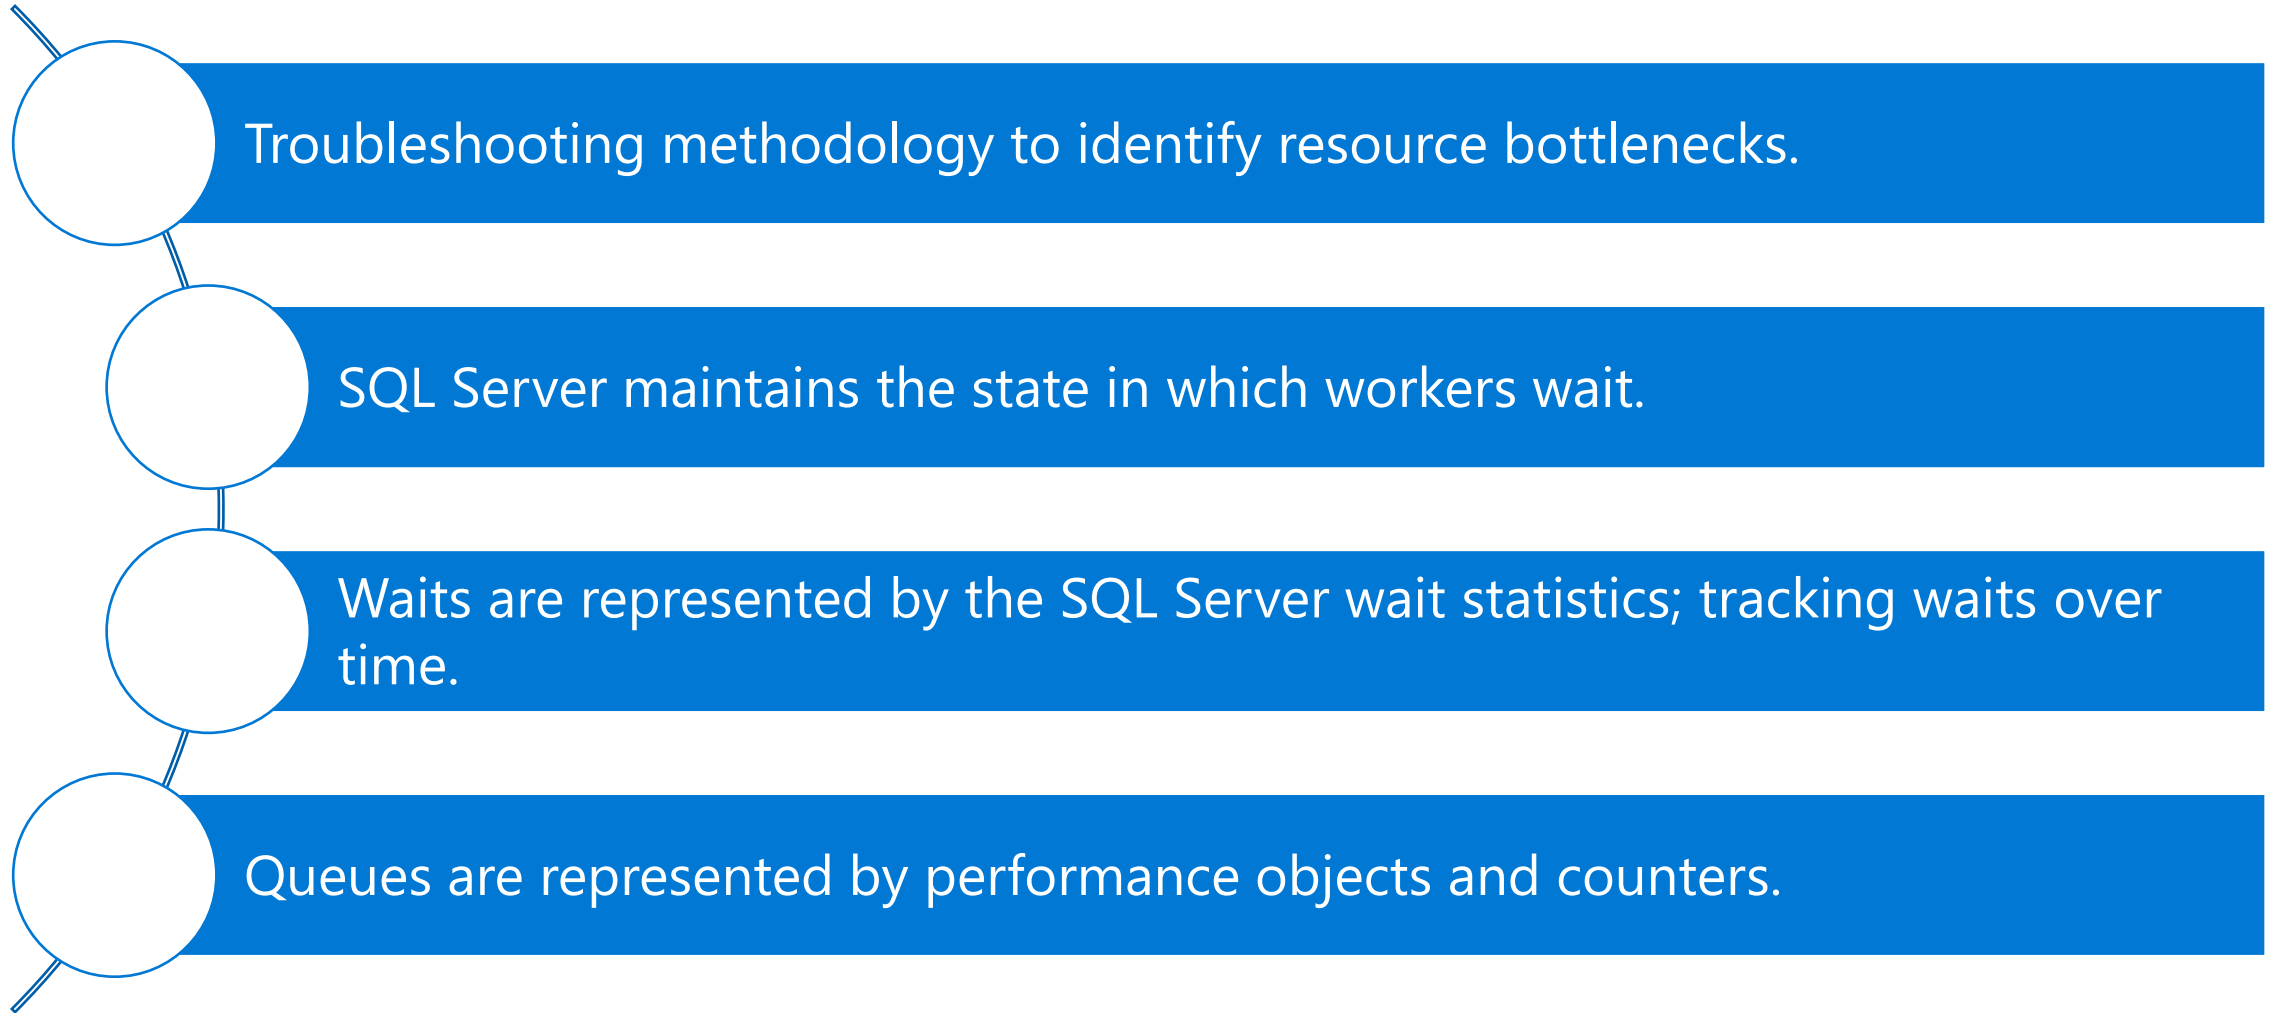
One SQLOS Scheduler
per core/logical
processor

Handles scheduling
tasks, I/O, and
synchronization of
resources

Work requests are
balanced across
schedulers based on
the number of active
tasks

Monitor using
`sys.dm_os_schedulers`

Waits and Queues



Thread States and Queues

Runnable: The thread is currently in the Runnable Queue waiting to execute. (First In, First Out).

Running: One active thread executing on a processor.

Suspended: Placed on a Waiter List waiting for a resource other than a processor. (No specific order).

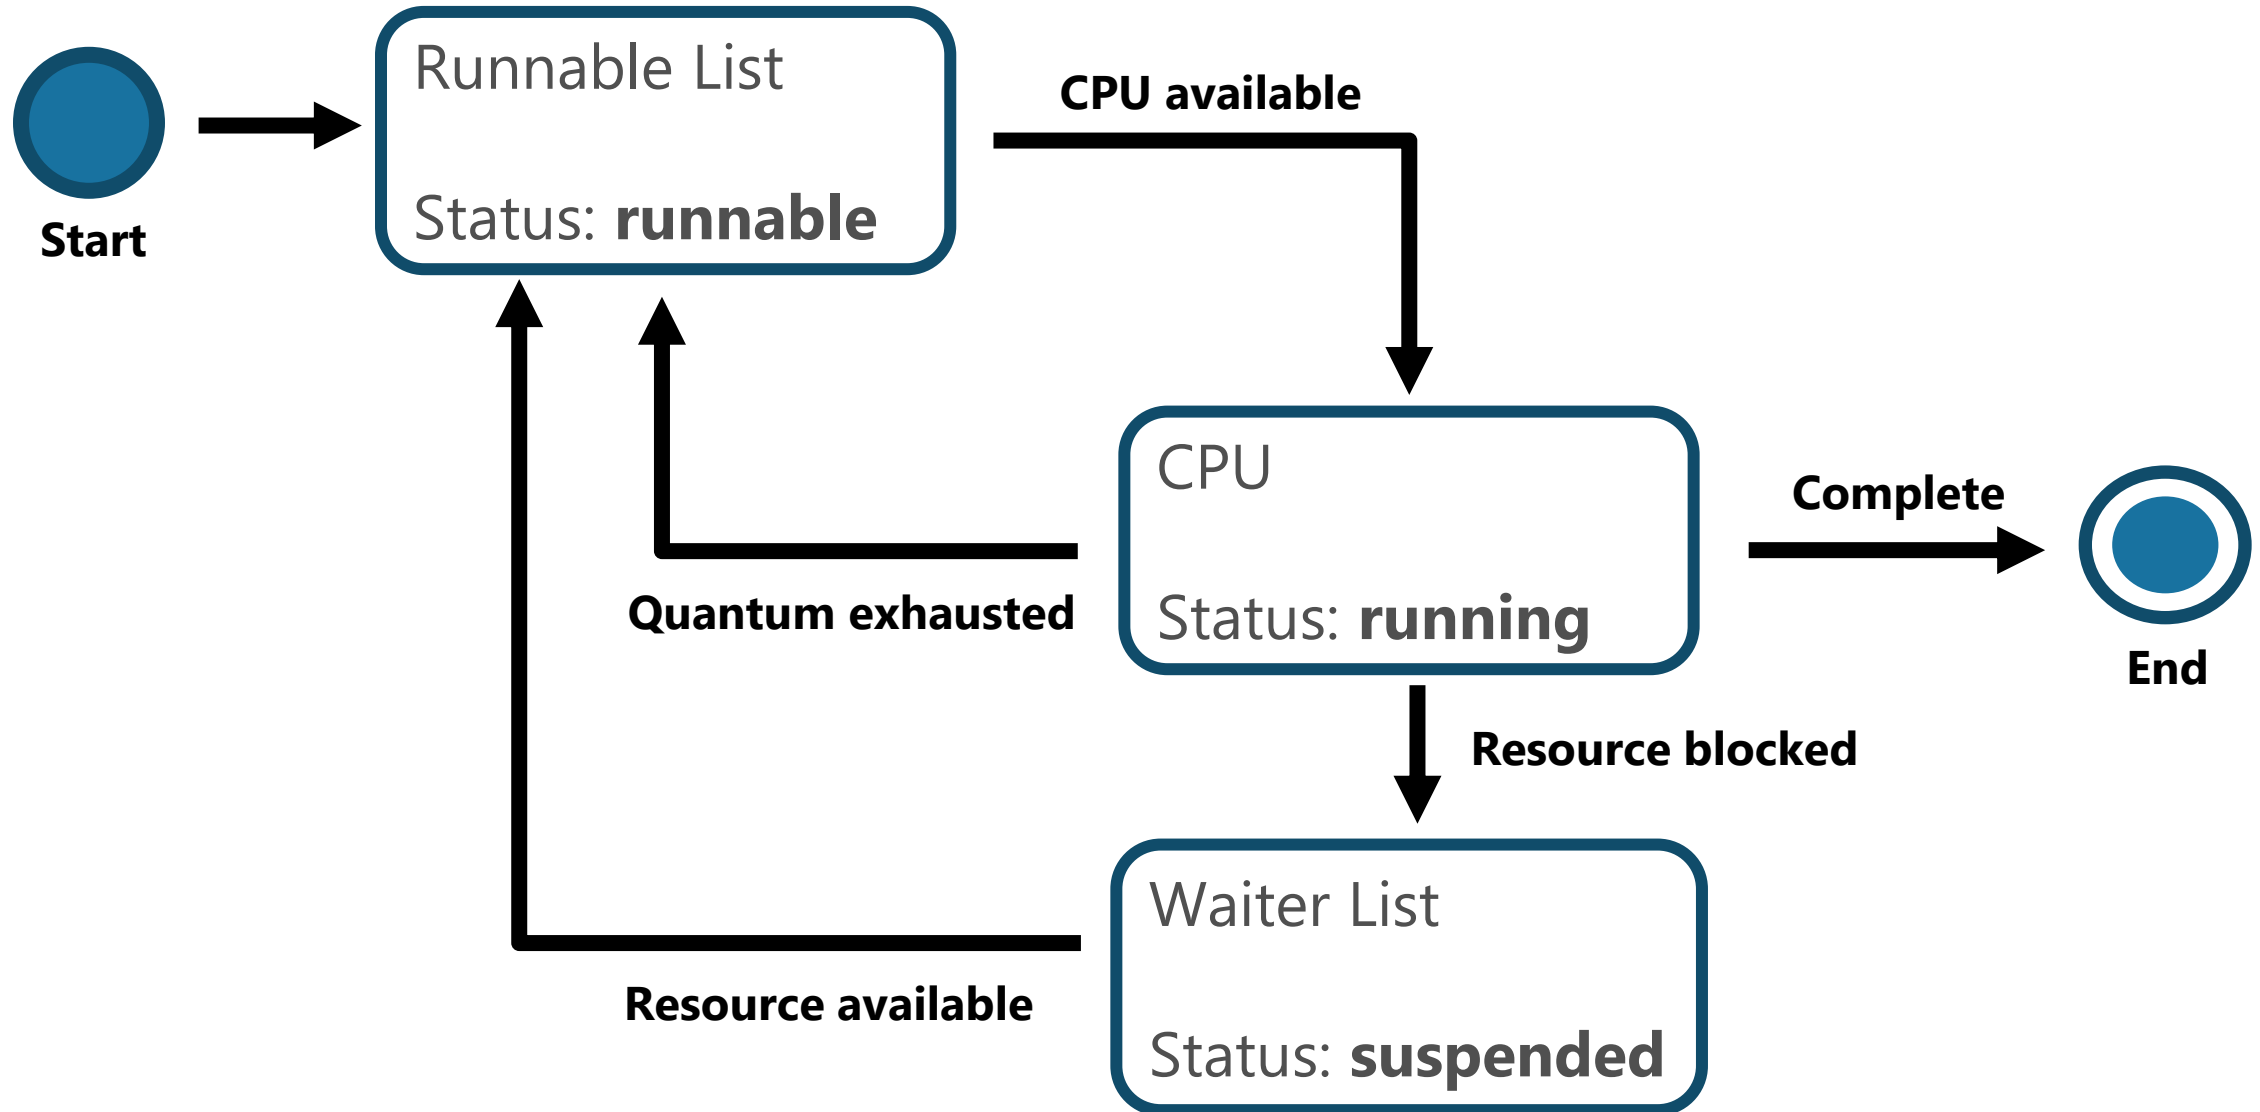
Yielding

In SQL Server, each thread is assigned a quantum (duration 4ms), with SQL Server using a cooperative model to ensure its CPU resources are shared amongst all the threads that are in a runnable state, preventing the 'starving' condition of any individual thread.

By design, a worker owns the scheduler until it yields to another worker on the same scheduler.

When no worker is currently on the Runnable list, the **yielding worker is allowed another quantum** or performs the necessary idle scheduler maintenance.

Yielding



Task Execution Model

Status: Running

session_id 51	Running
---------------	---------



Runnable Queue (Signal Waits)

Status: Runnable

session_id 51	Runnable
session_id 64	Runnable
session_id 87	Runnable
session_id 52	Runnable
session_id 56	Runnable

Wait Queue (Resource Waits)

Status: Suspended


session_id 73	LCK_M_S
session_id 59	NETWORKIO
session_id 56	Runnable
session_id 55	RESOURCE_SEMAPHORE
session_id 60	IO_Completion



SPID56 moved to the bottom of the Runnable queue.

Task Execution Model

- The full cycle between the several task states, for how many times it needs to cycle, is what we experience as the total query response time.



Relevant Dynamic Management Views (DMVs)

`sys.dm_os_wait_stats`

- Returns information about all the waits encountered by threads that ran.
- Includes wait type, number of tasks that waited in the specific wait type, total and max wait times, and the amount of signal waits.

`sys.dm_os_waiting_tasks`

- Returns information about the wait queue of tasks actively waiting on some resource.

`sys.dm_exec_requests`

- Returns information about each request that is in-flight.
- Includes session owning the request and status of the request, which will reflect the status of one or more tasks assigned to the request.

Waiting Tasks DMV

```
SELECT w.session_id, w.wait_duration_ms, w.wait_type,
       w.blocking_session_id, w.resource_description,
       s.program_name, t.text, t.dbid, s.cpu_time, s.memory_usage
FROM sys.dm_os_waiting_tasks as w
     INNER JOIN sys.dm_exec_sessions as s
         ON w.session_id = s.session_id
     INNER JOIN sys.dm_exec_requests as r
         ON s.session_id = r.session_id
     OUTER APPLY sys.dm_exec_sql_text (r.sql_handle) as t
WHERE s.is_user_process = 1;
```

session_id	wait_duration_ms	wait_type	blocking_session_id	resource_description
58	8563	LCK_M_S	62	keylock hobtid=72057594047365120 dbid=5 id=lock1...

Troubleshooting Wait Types

Aaron Bertrand – Top Wait Types

<https://sqlperformance.com/2018/10/sql-performance/top-wait-stats>

Paul Randal – SQL Skills Wait Types Library

<https://www.sqlskills.com/help/waits/>

