



SQL Server Indexes and Statistics

Module 5

Learning Units covered in this Module

- Lesson 1: Index Internals
- Lesson 2: Index Strategy
- Lesson 3: Index Monitoring and Fragmentation
- Lesson 4: SQL Server Statistics Internals

Lesson 1: Index Internals

Objectives

After completing this learning, you will be able to:

- Understand differences between heap and clustered objects.
- Describe B-Trees, and why they are beneficial.
- Understand the differences between clustered and non-clustered indexes.
- Describe the different methods by which data structures are accessed.



What is an Index?

An index is an on-disk structure associated with a table that speeds retrieval of rows.


An index contains keys built from one or more columns in the table.



What is SARGability?

A SARGable item in a search predicate is able to use an index.

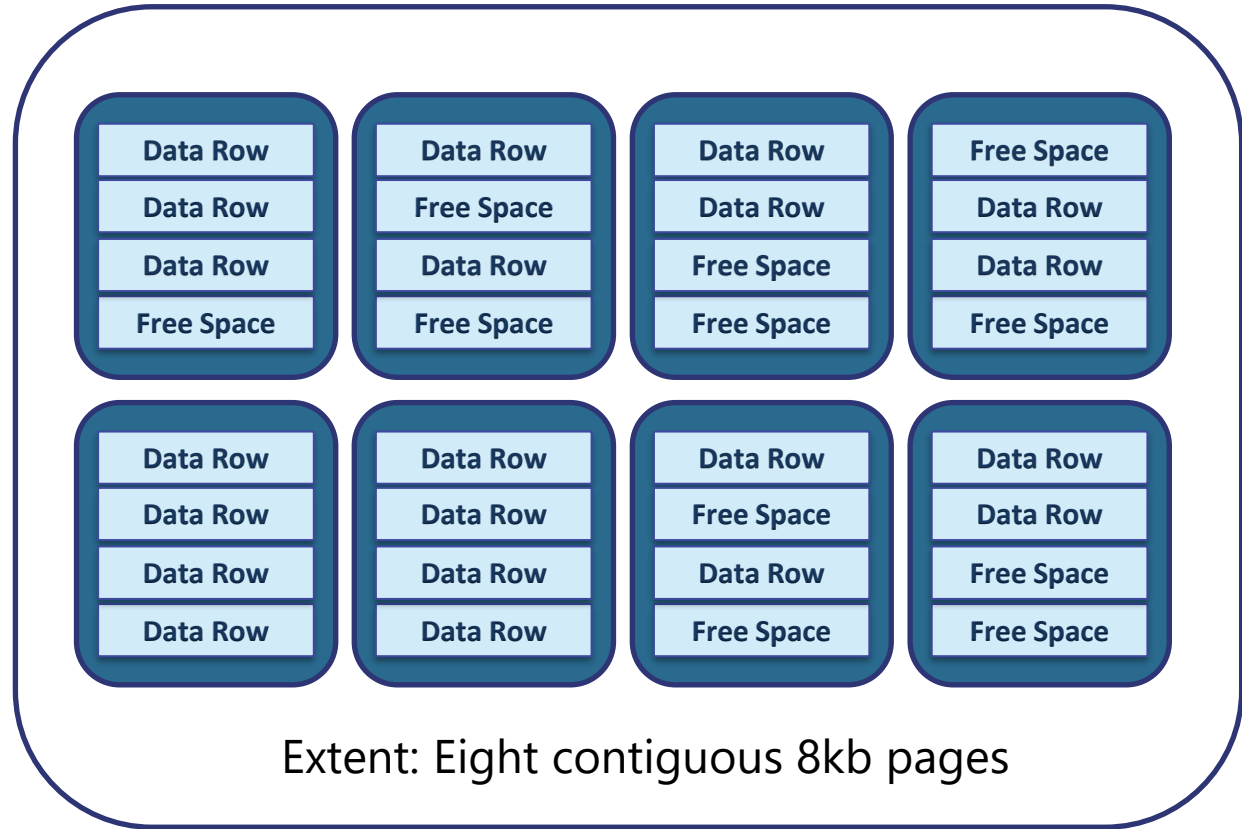
Non-SARGable expressions can significantly slow down queries.



How Data is stored in a Database



Primary Data File (.mdf)
Secondary Data File (.ndf)



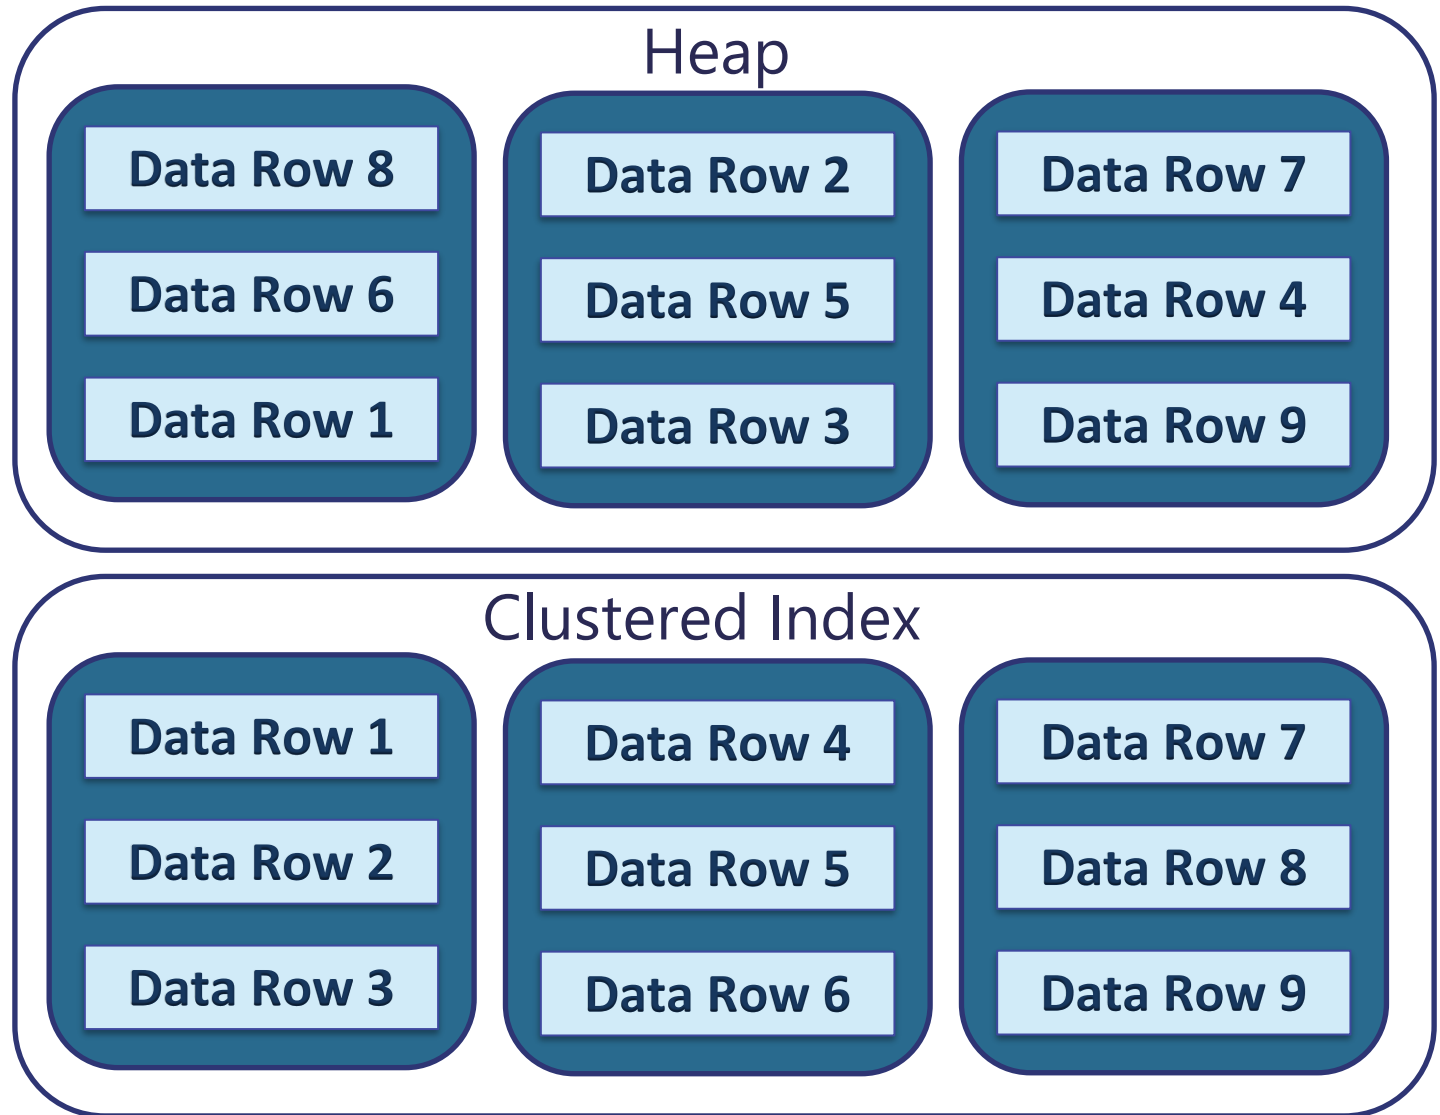
Uniform extents: Pages used by a single object.

Mixed extents: Pages used by different objects.

How Data is Stored in Data Pages

Data stored in a Heap is not stored in any order and normally does not have a Primary Key.

Clustered Index data is stored in sorted order by the Clustering key. In many cases, this is the same value as the Primary Key.



Heap

A heap is a table without a clustered index

Unordered masses of data

Can be good for quickly importing large sets of data

Not a good idea for reporting-based data structures


Use the ALTER TABLE...REBUILD command to "rebuild" a heap table

Do not use a heap

- When the data is frequently returned in a sorted order.
- When the data is frequently grouped together.
- When ranges of data are frequently queried from the table.
- When there are no nonclustered indexes and the table is large.

Heap Structures

Heaps have one row in `sys.partitions`, with `index_id = 0` for each partition used by the heap



The screenshot shows a SQL query window with the following text:

```
1 SELECT * FROM Membership
2 WHERE FirstName = 'Janice'
3
```

Below the query, the "Execution plan" tab is selected. It displays the following information:

Query 1: Query cost (relative to the batch): 100%

SELECT * FROM [Membership] WHERE [FirstName]=@1

The execution plan consists of three operators:

- SELECT**: Cost: 0 %
- Compute Scalar**: Cost: 0 %
- Table Scan [Membership]**: Cost: 100 %, 0.000s, 1 of 1 (100%)

The operators are connected by lines, indicating the flow of data from the Table Scan through the Compute Scalar operator to the SELECT operator.

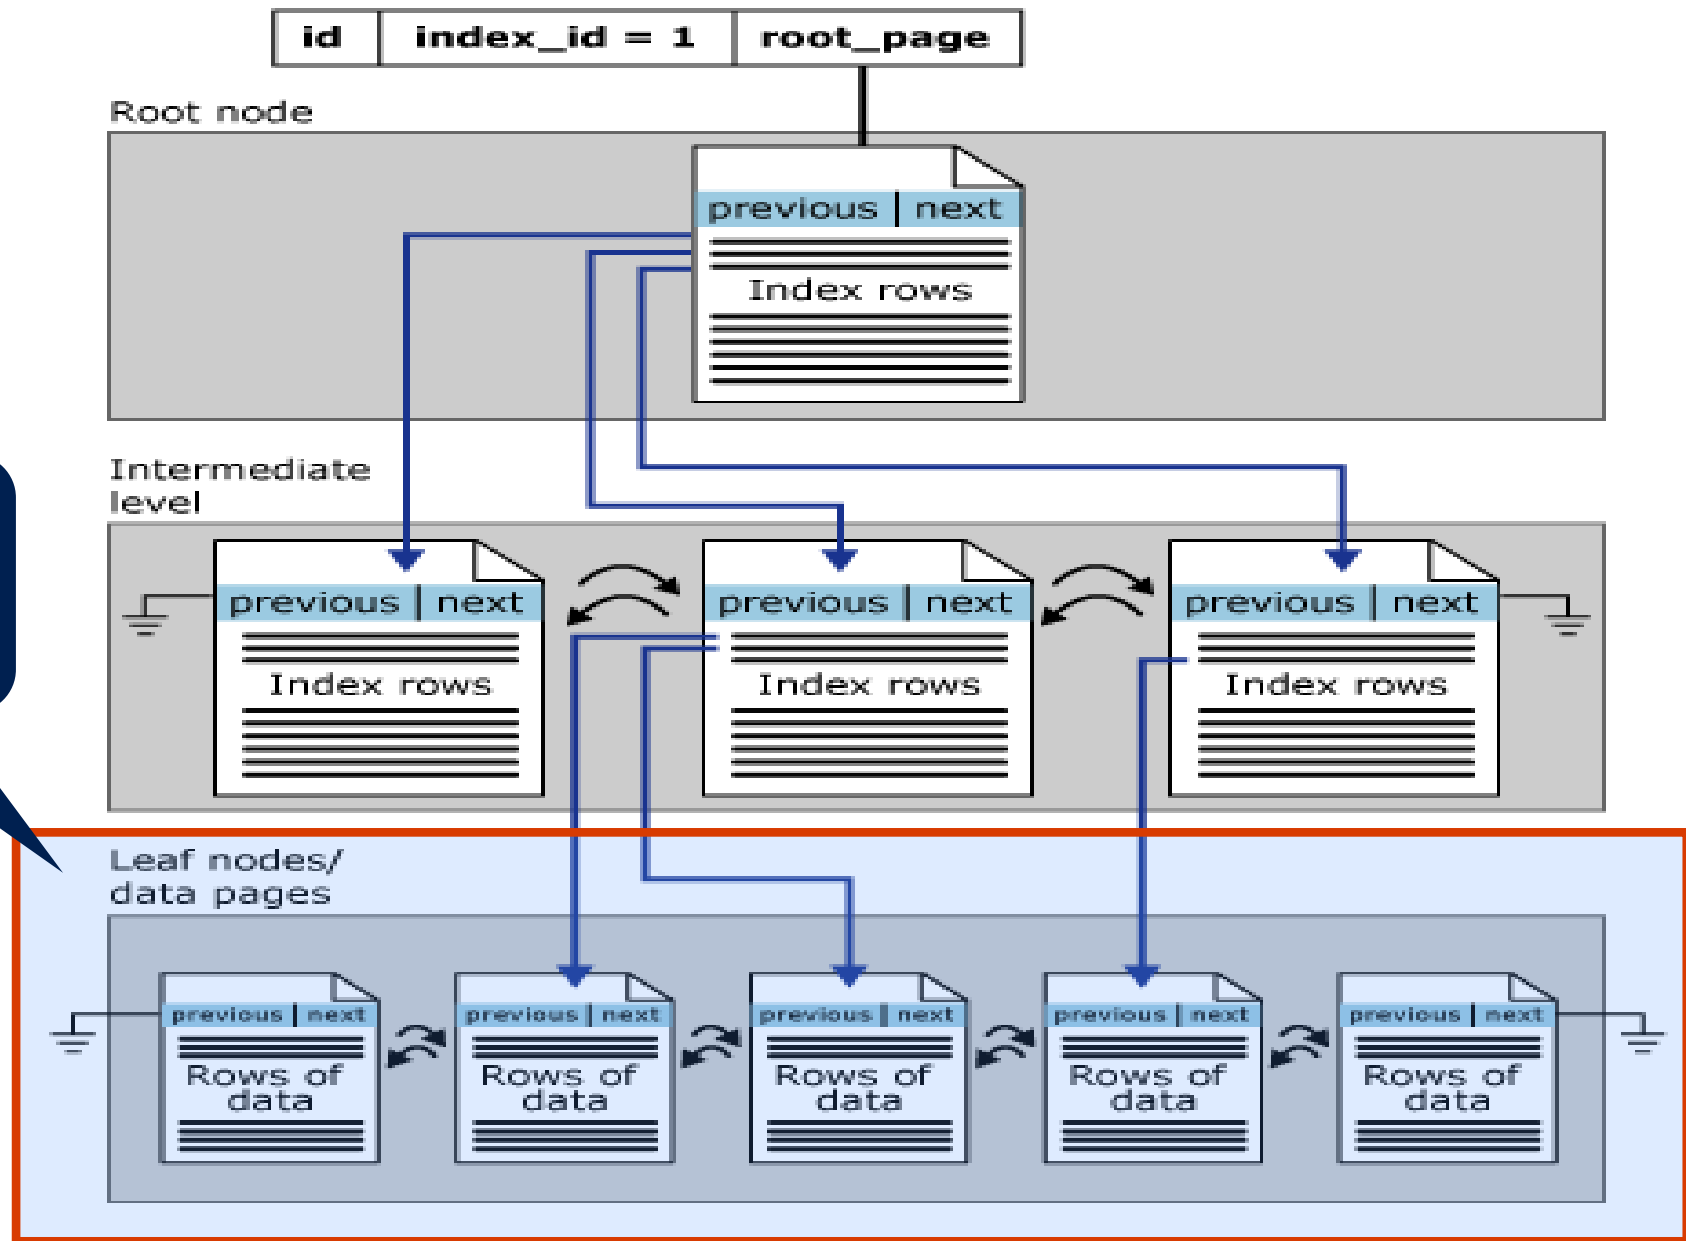
Clustered Indexes

An ordered data structure that is implemented as a Balanced (B) Tree.

When a table is clustered, the leaf level of the index contains ALL data in the table. This means that the clustered index IS the table. This is also why there is only one per table.

The leaf level of the clustered index contains data pages.

Clustered Index Structure



**Leaf Node
Contains
Data Pages**

Demonstration

Examining index metadata and Building a B-Tree

- This demonstration reviews index, partition, and allocation unit metadata.
- Show how a B-Tree begins to form as rows are added to a table.



Non-Clustered Indexes

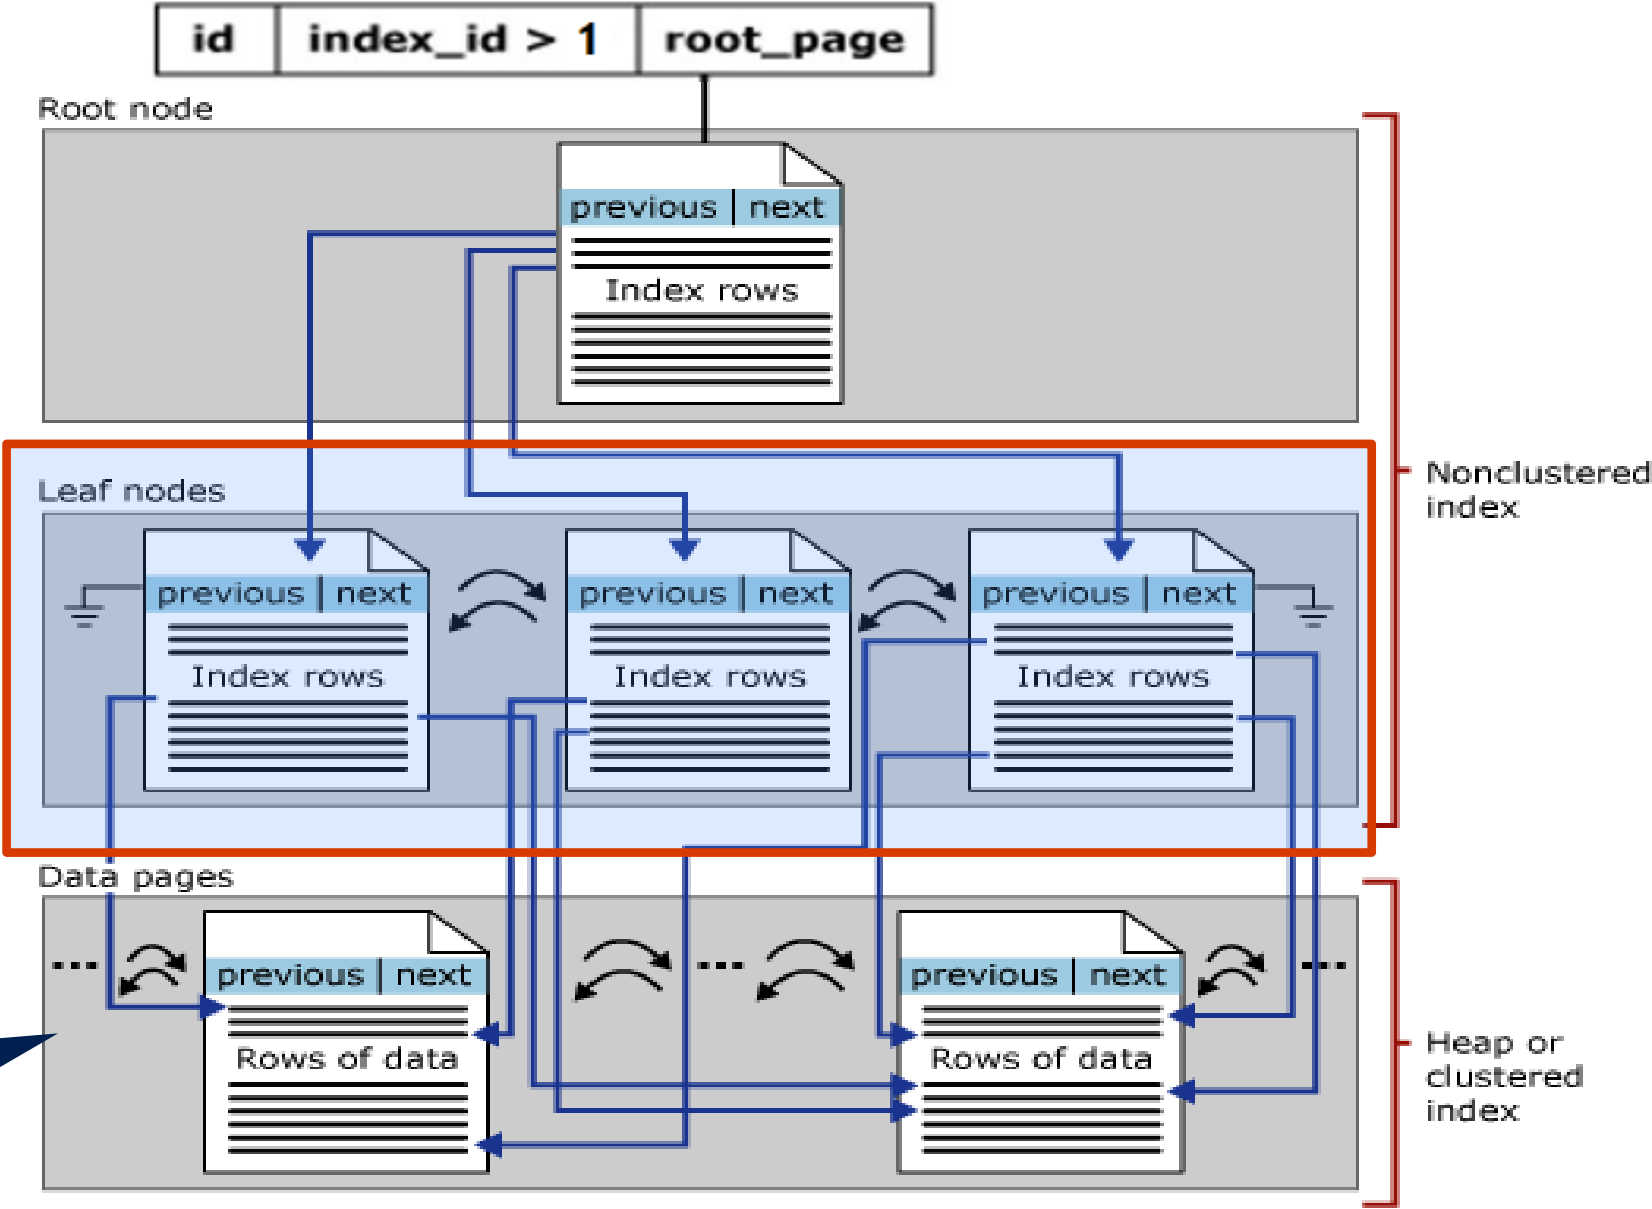
Same B-tree data structure as a clustered index.

It is a separate structure built on top of a Heap or Clustered Index.

Only contains a subset of the columns in the base table

The leaf level contains only the columns defined in the index as well as the clustered key/heap ID that points to the base table structure.

Non-Clustered Index Structure



**Leaf Node
has Pointers
to Base Table**

**Base Table
(Heap or
Clustered
Index)**

Clustered vs Nonclustered Indexes

An index is an on-disk structure associated with a table or view that speeds retrieval of rows.

Clustered Indexes

- Defines the order in which data is physically stored in a table.
- Table data can be sorted in only one way.
- Leaf level has data rows stored with index.
- When a table has a clustered index, the object is called clustered table.
- Cluster key is added to nonclustered index (as the pointer), keep it as narrow as possible.

Non-clustered Indexes

- Separate structure from base table.
- Contains a pointer back to base table called:
 - Row ID (RID when base table is HEAP)
 - Key (KEY when base table is Clustered)
- "Skinny" data structure as it contains a subset of base table only.
- To by-pass index key limits (1,700 bytes), non-key columns can be added to leaf level.
- As Leaf level contains fewer columns than base table, the non-clustered index uses fewer pages than the corresponding base table.

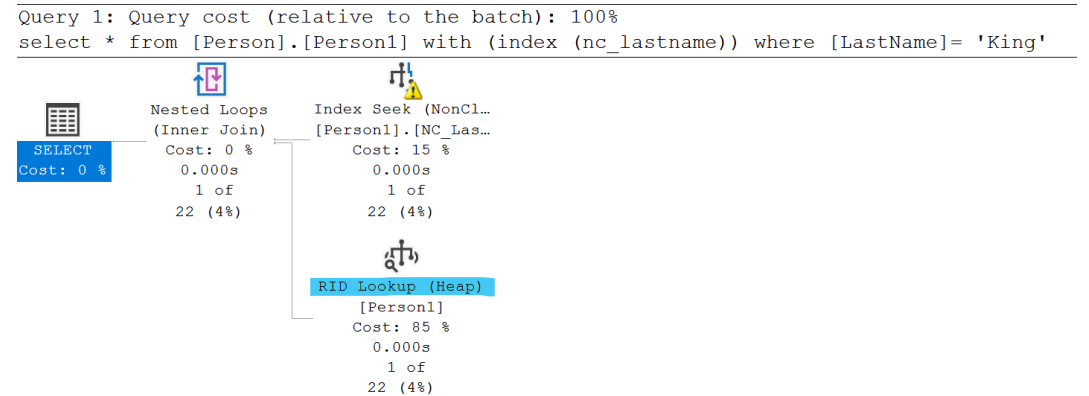

Bookmark Lookup or RID lookup

Occurs when SQL uses a nonclustered index to satisfy all or some of a query's predicates, but it doesn't contain all the information needed to cover the query.

Lookup effectively join the nonclustered index back to the clustered index or heap.

- If table has clustered index, it is called **bookmark lookup** (or key lookup)
- If the table is a heap with a non-clustered index, it is called **RID lookup**

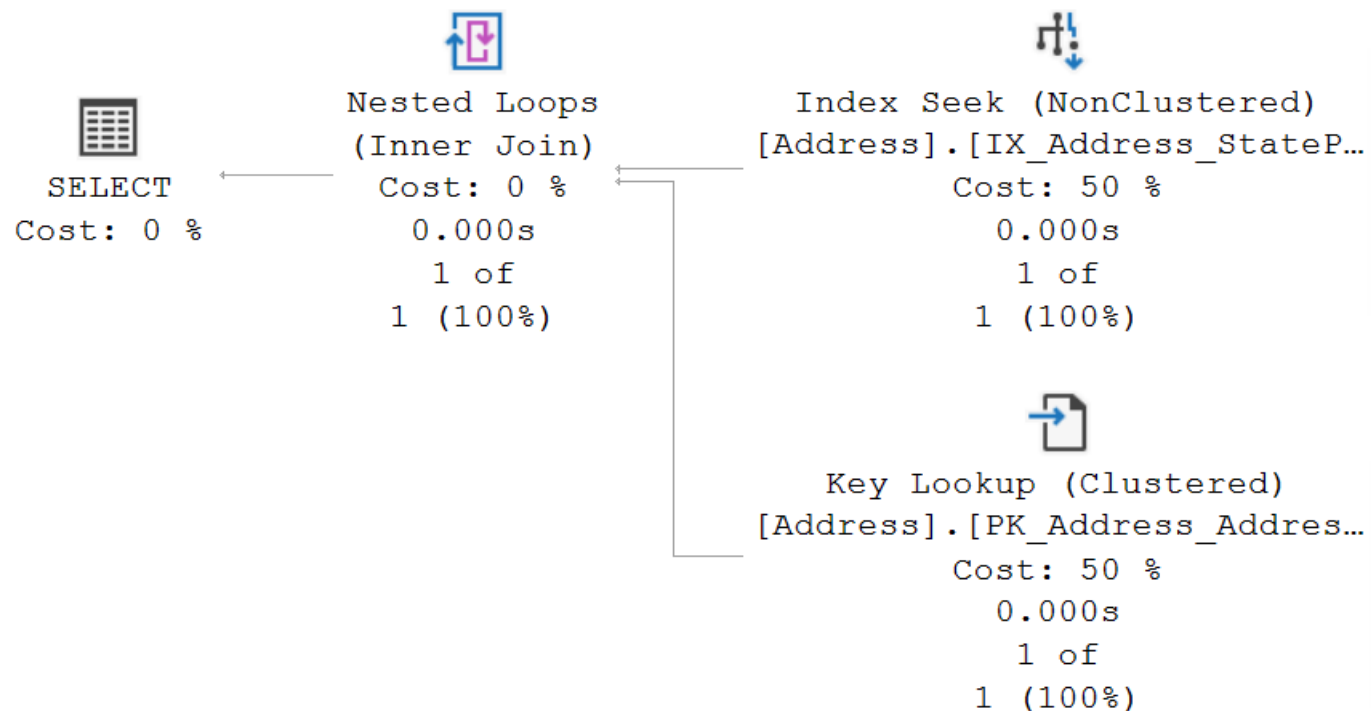
This is an expensive operation.



Key Lookup

Query 1: Query cost (relative to the batch): 100%

```
SELECT [AddressID],[StateProvinceID],[City] FROM [Person].[Address] WHERE [StateProvinceID]=@1
```



Object

[AdventureWorks2016].[Person].[Address].
[IX_Address_StateProvinceID]

Output List

[AdventureWorks2016].[Person].[Address].AddressID,
[AdventureWorks2016].[Person].[Address].StateProvinceID

Object

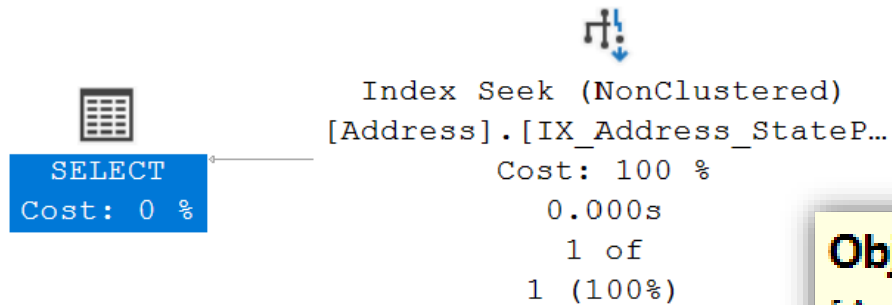
[AdventureWorks2016].[Person].[Address].
[PK_Address_AddressID]

Output List

[AdventureWorks2016].[Person].[Address].City

Non-Clustered Index with Included Column

Query 1: Query cost (relative to the batch): 100%
SELECT [AddressID],[StateProvinceID],[City] FROM [Person].[Address] WHERE [StateProvinceID]=@1



Object

[AdventureWorks2016].[Person].[Address].
[IX_Address_StateProvinceID]

Output List

[AdventureWorks2016].[Person].[Address].AddressID,
[AdventureWorks2016].[Person].[Address].City,
[AdventureWorks2016].[Person].[Address].StateProvinceID

Demonstration

Access Methods

Demonstrate nonclustered indexes with included columns



Questions?




Lesson 2: Index Strategy

Objectives

After completing this learning, you will be able to:

- Describe best practices associated with clustered index design.
- Consider factors involved in designing a nonclustered indexing strategy.
- Explain how an indexing strategy can reduce logical reads.
- Describe use of Filter indexes and Indexed views.
- Explain the role of Fill factor.
- Use a technique to reduce the fragmentation.
- Apply Index Compression.



Characteristics of a Good Clustering Key

Narrow

- Use a data type with a small number of bytes to conserve space in tables and indexes

Unique

- To avoid SQL adding a 4-byte uniquifier

Static

- Allows data to stay constant without constant changes which could lead to page splits

Increasing

- Allows better write performance and reduces fragmentation issues

Best Practices for Clustered indexes

The primary goal of any indexing strategy is to make queries faster

Best Practice	Seeks, not scans	Avoid key lookups	Get most rows on the index page	Get most rows on the data page	Avoid page splits
Use frequently searched value	X	X			
Use narrow keys			X		
Use unique values			X	X	
Use static values					X
Use ascending or descending key					X

Clustered Index vs Primary key considerations

Frequently, clustering on the primary key supports the clustered index best practices

- Query predicate use Primary Key
- Primary keys are involved in joins to foreign keys
- Primary keys are unique
- Primary keys are typically ascending
- Primary keys are typically not updated

Do not cluster on the Primary key

- When table is accessed by other keys than the primary key
- Common issue accessing association or child table, when these tables require different primary key than parent table
- Use `sys.dm_db_index_usage_stats` to review data access patterns of index

Non-Clustered Index Strategy

Index columns used in WHERE clauses and JOINs.

Keep indexes to a minimum, to minimize impact on DML and log writes.

Two basic approaches: Single column indexes or Multicolumn indexes

Easy to create too many nonclustered indexes.

Specific indexes are appropriate for high-impact queries.

Ideally the focus should be writing queries to use existing indexes, rather than on adding more indexes.

Single Vs Multi-column indexes

Single-column Indexing

- Good choice for columns that are highly selective, or columns referenced often by criteria.
- When filtering data from table with multiple single column indexes, SQL uses the index with higher cardinality column, it may be better to not index the low cardinality column.

Multi-Column Indexing

- Good choice for low cardinality columns.
- Highest cardinality field first, to zoom into a narrow range of the B-Tree as quickly as possible.
- Equality columns before inequality columns.

Multi-Column Indexing Access

Seek only happens if you search for the columns in the specified order of how the B-tree is organized.

```
CREATE INDEX IX1 ON TABLE (StateID, City, PostalCode)
```

Effect of column in different WHERE clause.


- `WHERE StateID = 79` -- seek
- `WHERE StateID = 79 AND City = Bothell` -- seek
- `WHERE StateID = 79 AND PostalCode = 98011` -- seek then scan
- `WHERE PostalCode = 98011` -- scan

Lesson 3: Index Monitoring and Fragmentation

Objectives

After completing this learning, you will be able to:

- Monitor index usage.
- Describe page splits.
- Understand problems caused by page splits.
- Monitor index fragmentation and free space.
- Review options to remove fragmentation and free space.



Why to monitor Index usage

Data access and distribution changes over time.

By reviewing index usage, we can ensure that the indexes are appropriate for the current workload.

Use `sys.dm_db_index_usage_stats` and `sys.dm_db_index_operational_stats` to monitor index usage.

Database standard report also monitors Index usage and Physical Statistics.

Clustered Index usage patterns

Pattern	Interpretation
Seeks, but no scans	<ul style="list-style-type: none">• May not need maintenance. Seeks are not impacted by fragmentation
Clustered index with High number of scans	<ul style="list-style-type: none">• Review missing index recommendations.• Consider changing Clustering key.• Simplify queries with complex WHERE clauses
Clustered index with Low number of seeks and high number of lookups	<ul style="list-style-type: none">• May need to change Clustering Key• Look how often bookmark lookups occur.

Non-Clustered Index usage patterns

Pattern	Interpretation
Seeks, but no scans	<ul style="list-style-type: none">• May not need maintenance.• Seeks are not impacted by fragmentation
Non-Clustered index with High number of scans	<ul style="list-style-type: none">• Consider compression.• Check for non-SARGable queries.• Review missing index recommendations
Index with No/low usage	<ul style="list-style-type: none">• May be able to drop or disable index
Non-Clustered Index Used at specific times only	<ul style="list-style-type: none">• May be able to drop index, recreate when needed

sys.dm_db_index_usage_stats

- Tracks seeks, scans, and lookups
- Provides date of last access
- Counters are incremented once per query execution
- User and system access (such as index rebuilds) are tallied separately
- Counters are initialized to at SQL Server (MSSQLSERVER) service is restarts

```
SELECT
    Distinct OBJECT_NAME(Indx.OBJECT_ID) Table_Name ,Indx.name IndexName
    ,Indx.type_desc IndexType,Iusg.user_seeks Seeks,Iusg.user_scans Scans
    ,Iusg.user_lookups Lookups ,Iusg.user_updates Updates ,Iusg.last_user_seek LastSeek
    ,Iusg.last_user_scan LastScan,Iusg.last_user_lookup LastLookup ,Iusg.last_user_update LastUpdate
FROM
    SYS.INDEXES Indx
    INNER JOIN SYS.DM_DB_INDEX_USAGE_STATS Iusg
        ON Iusg.index_id = Indx.index_id AND Iusg.OBJECT_ID = Indx.OBJECT_ID
    INNER JOIN SYS.DM_DB_PARTITION_STATS PrtSts
        ON PrtSts.object_id=Indx.object_id
WHERE
    OBJECTPROPERTY(Indx.OBJECT_ID,'IsUserTable') = 1
```

sys.dm_db_index_operational_stats

- Returns information about the lower-level I/O activities
- Takes database_id, the object_id, the index_id and the partition_number as parameters
- Memory-optimized indexes do not appear in this DMV

```
SELECT
    *
FROM
    SYS.DM_DB_INDEX_OPERATIONAL_STATS (NULL,NULL,NULL,NULL ) opsts
    INNER JOIN SYS.INDEXES AS Indx
        ON Indx.[OBJECT_ID] = opsts.[OBJECT_ID]
        AND Indx.INDEX_ID = opsts.INDEX_ID
WHERE
    OBJECTPROPERTY(opsts.[OBJECT_ID], 'IsUserTable') = 1
```

sys.dm_db_index_physical_stats

- Returns size and fragmentation information for the data and indexes of the specified table or view in SQL Server.
- Takes database_id, the object_id, the index_id and the partition_number as parameters
- Memory-optimized indexes do not appear in this DMV

```
SELECT OBJECT_SCHEMA_NAME(I.object_id) AS SchemaName, OBJECT_NAME(I.object_id) AS  
TableName, I.name, I.Index_ID, IPS.partition_number, IPS.avg_fragmentation_in_percent,  
IPS.page_count, IPS.avg_page_space_used_in_percent  
FROM sys.indexes as I  
INNER JOIN sys.dm_db_index_physical_stats(DB_ID(), NULL, NULL, NULL, 'Limited') as IPS  
ON I.object_id = IPS.object_id  
AND I.index_id = IPS.index_id  
--WHERE IPS.avg_fragmentation_in_percent >30  
--AND IPS.page_count >1000
```

Fragmentation

A fragmented table/Index is when some of its data pages point to pages that are not in sequence.



Logical fragmentation

- Occurs when leaf level pages are not physically corresponding to the logical order of the index:
 - Pages are not in the most efficient order for scanning purposes.
- Limits the efficiency of read-ahead scans, but not seeks.




Page density


- How full a page is when a rebuild/reorganization occurs.
- The fuller a page is, the more likely page splits occur when data is modified.
- The less full a page is, the more wasted space in the buffer pool when reading pages.

Page Splits

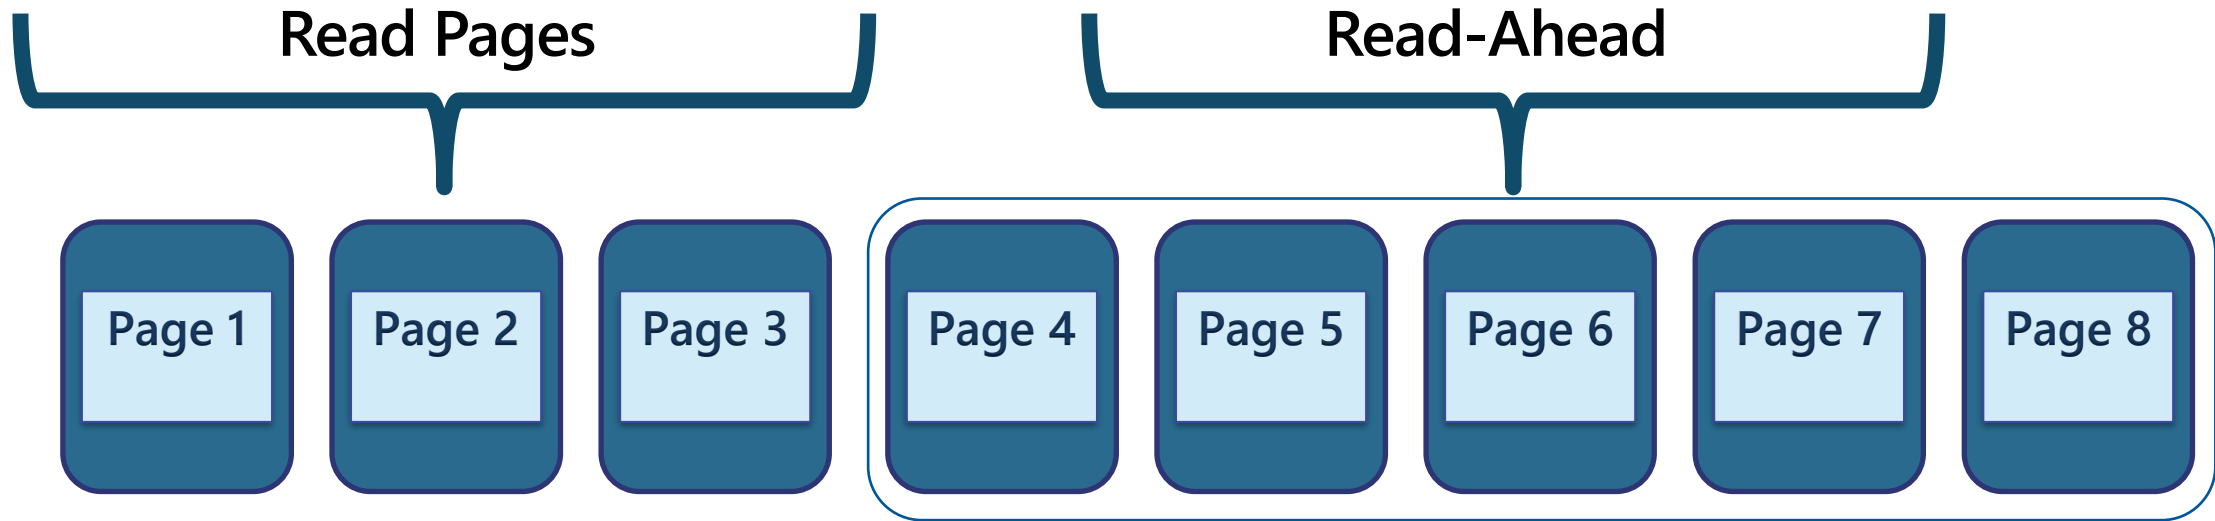
Inserting a new record, causing a page split



After the insert we see disproportionately greater overhead



Read-Ahead Scans



- Read-ahead anticipates the data and index pages needed to fulfill a query execution plan and brings the pages into the buffer cache before they are used by the query.
- The read-ahead mechanism allows the Database Engine to read up to 64 contiguous pages (512KB) from one file.

SQL Index Fragmentation and Maintenance options

- Maintenance plans
- SSMS -> Index Physical Statistics report
 - Will report fragmentation and recommend solution at database level
- Custom Solution
 - Find fragmentation using `sys.dm_db_index_physical_stats`
 - Correct fragmentation
- Can remove Index fragmentation using Rebuild or Reorganize option

	ALTER INDEX... REBUILD	ALTER INDEX ... REORGANIZE
Removes fragmentation	X	X
Removes free space	X	
Resets fill factor	X	
Updates statistics	X	

Fill Factor

Can address performance issues due to fragmentation.

Use When

- Specifying the amount of free space on a data or index page.
- Reducing logical fragmentation.

Do not use when

- Low fragmentation
- High seeks and low scans
- Index is not scanned

```
ALTER INDEX [pk_bigProduct]
ON [dbo].[bigProduct]
REBUILD WITH (PAD_INDEX = ON, FILLFACTOR = 90)
```

Reducing fragmentation

Reducing Fragmentation

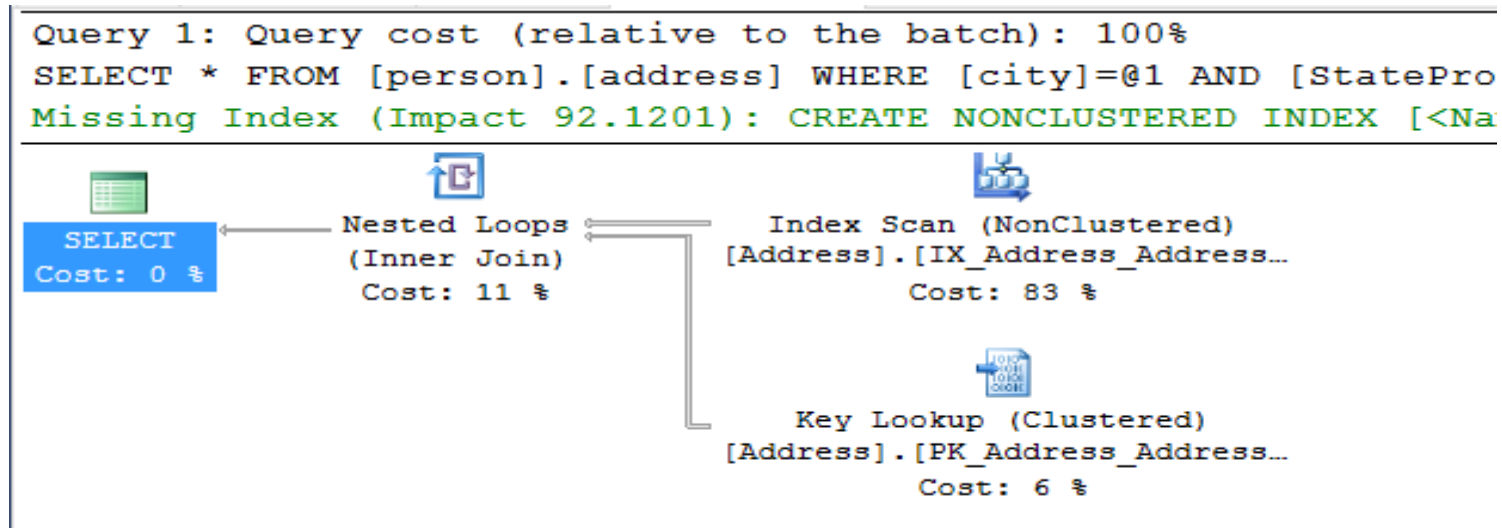
- Using an ascending key (not always possible).
- Using an appropriate fill factor for the workload.
- Update in sets, rather than one at a time.
- Do not insert with immediate update.

Missing Index Recommendations

SQL Server identifies indexes that would help a query's performance

The recommendation is included in the query plan

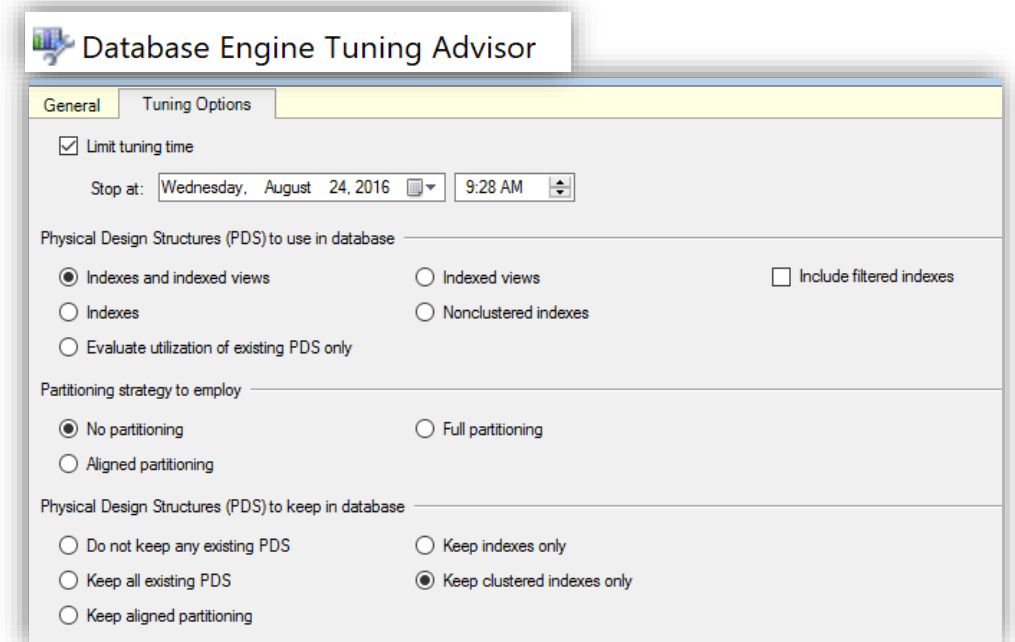
The cost savings are aggregated in DMVs to help identify the most beneficial indexes



Choosing Indexes

Following sources can help you choosing indexes

- Missing Index recommendations
- `sys.dm_db_index_usage_stats`
- Database Tuning Advisor



Demonstration

Analyzing Index Usage

Show how index usage patterns can help identify index and query improvements.



Lesson 4: SQL Server Statistics Internals

Objectives

After completing this learning, you will be able to:

- Understand statistics, and to retrieve their contents.
- Review database options used to control statistics creation and update.
- Use different methods to read statistics information.



SQL Server statistics

What are the statistics?

Statistics contain statistical information about distribution of values in one or more columns of a table or index.

It is stored as binary large objects (BLOBs).

It is used by the Query Optimizer (QO) to estimate the *cardinality*, or number of rows, in the query result, and enable the creation of high-quality query plans.

Statistics are created:

- Intrinsically when indexes are created.
- Manually by using CREATE STATISTICS command.
- Automatically, to support WHERE clauses (if AUTO_CREATE_STATISTICS is ON).

Statistics components

Density Vector

Density is information about the number of duplicates in each column or combination of columns

- Used when query predicate contains variable

WHERE col = @variable

or when a stored procedure uses query on a modified parameter:

WHERE col = @local_variable

Histogram

A **histogram** measures the frequency of occurrence for each distinct value in a data set.

- Used when query predicate contains

WHERE col = 'literal'

or when a stored procedure uses a query on a parameter

WHERE col = @parameter


Statistics components

Histogram

It is computed from the values in the leftmost key column of the statistics object.

Data is selected from table or view in one of two ways:

- Statistically sampling rows.
Data shown contains estimates of rows and distinct values.
- Performing a full scan of all rows.
Data shown contains actual numbers.



Showing Statistics

```
DBCC SHOW_STATISTICS ('Sales.SalesOrderDetail', 'IX_SalesOrderDetail_ProductID')  
WITH STAT_HEADER, HISTOGRAM
```

Results									
Messages									
Name	Updated	Rows	Rows Sampled	Steps	Density	Average key length	String Index	File	
IX_SalesOrderDetail_ProductID	Nov 7 2012 6:44PM	121317	121317	200	0.0078125	12	NO		N
RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS					
730	0	288	0	1					
732	0	130	0	1					
738	154	600	2	77					
741	167	94	1	167					
742	0	288	0	1					

```
SELECT  
    ProductID, RecordCount = COUNT(*)  
FROM Sales.SalesOrderDetail  
WHERE  
    ProductID >= 732 AND  
    ProductID <= 738  
GROUP BY ProductID
```

Results	
Messages	
ProductID	RecordCount
732	130
733	44
736	110
738	600

Showing Statistics

Statistics Properties - IX_SalesOrderDetail_ProductID

Select a page

General

Details

Filter

Connection

Server:
jdsq1-one.database.windows.net

Connection:
johndeardurff

View connection properties

Progress

Ready

Script Help

Table Name:
SalesLT.SalesOrderDetail

Statistics Name:
IX_SalesOrderDetail_ProductID

Statistics for INDEX 'IX_SalesOrderDetail_ProductID'.

Name	Updated
IX_SalesOrderDetail_ProductID	Aug 20 2019 1:09PM
All Density	Average Length
0.007042253	4
0.001845018	8
0.001845018	12ProductID,
Histogram Steps	
RANGE_HI_KEY	RANGE_ROWS
707	0
708	0
711	0
712	0
714	0
715	0
716	0
718	2
722	0
738	0
739	0
742	0
743	0
747	0
748	0

OK

Cancel

Automatically created statistics

Database options that affects automatic statistics creation and update

AUTO_CREATE_STATISTICS

AUTO_UPDATE_STATISTICS

AUTO_UPDATE_STATISTICS_ASYNC

INCREMENTAL

Use the defaults unless you NEED to do otherwise.

Do not enable auto-create statistics on SharePoint content database.

The small sample rate of AUTO_UPDATE_STATISTICS can cause some workloads to choose sub-optimal execution plans.

Manually created Statistics

For most queries, the query optimizer generates necessary statistics for a high-quality query plan.

In a few cases, additional statistics is needed to improve query performance.

```
CREATE STATISTICS ContactPromotion1  
ON Person.Person (BusinessEntityID, LastName, EmailPromotion)
```

Available tools to review statistics

	Metadata	Last Update Date	Sampling Rate	Row Mod Counter	Density Vector	Histogram
sys.stats	X					
STATS_DATE()		X				
sys.dm_db_stats_properties (object_id, stats_id)		X	X	X		
sys.dm_db_incremental_stats_properties (object_id, stats_id)		X	X	X		
sys.dm_db_stats_histogram (object_id, stats_id)						X
SQL Server Mgt Studio	X	X	X		X	X
DBCC SHOW_STATISTICS		X	X	X	X	X

Demonstration

Exploring statistics



Questions?



